



NEW YORK UNIVERSITY

Convolutional Networks

<http://bit.ly/DLSP20>

Yann LeCun

NYU - Courant Institute & Center for Data Science

Facebook AI Research

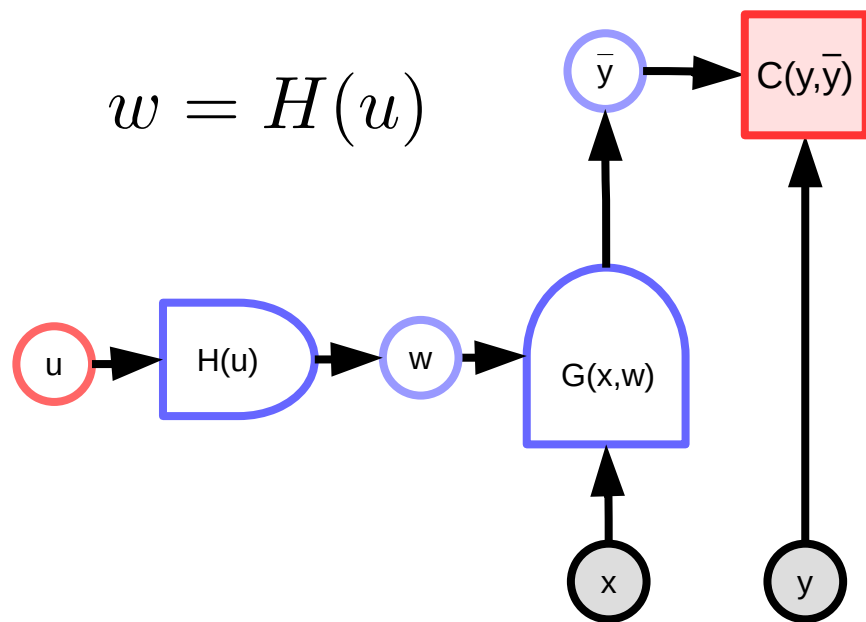
<http://yann.lecun.com>

TAs: Alfredo Canziani, Mark Goldstein

Deep Learning, NYU, Spring 2020

Parameter transformations

- When the parameter vector is the output of a function

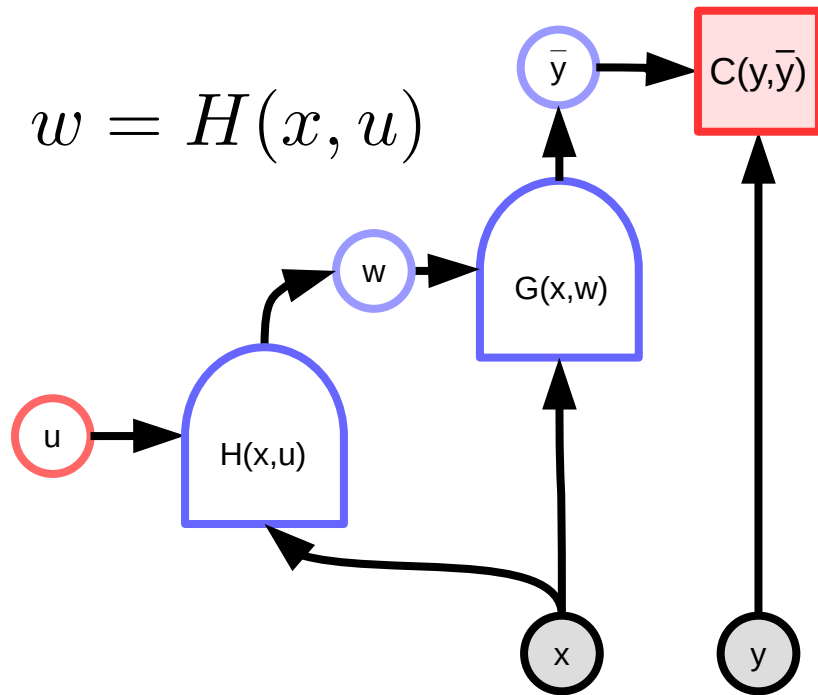


$$u \leftarrow u - \eta \frac{\partial H^T}{\partial u} \frac{\partial C^T}{\partial w}$$

$$w \leftarrow w - \eta \frac{\partial H}{\partial u} \frac{\partial H^T}{\partial u} \frac{\partial C^T}{\partial w}$$

$$[N_w \times N_u] \quad [N_u \times N_w] \quad [N_w \times 1]$$

“Hypernetwork”



- ▶ When the parameter vector is the output of another network $H(x, u)$
- ▶ The weights of network $G(x, w)$ are dynamically configured by network $H(x, u)$
- ▶ The concept is very powerful
 - ▶ The idea is very old

Simple parameter transform: weight sharing

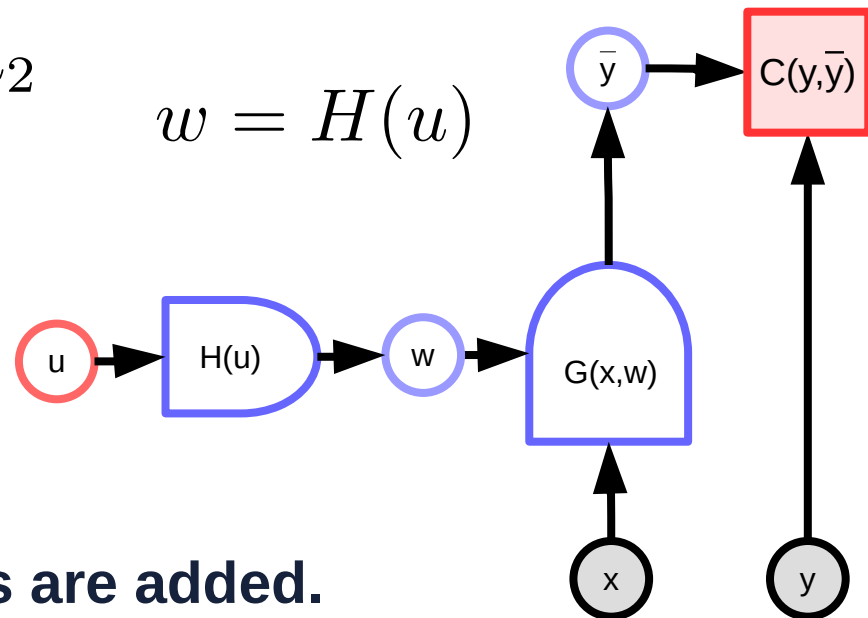
- ▶ Function $H(u)$ replicates one component of u into multiple components of w

$$w_1, w_2 = u_1 \quad w_3 = w_4 = u_2$$

- ▶ H is like a “Y” branch.

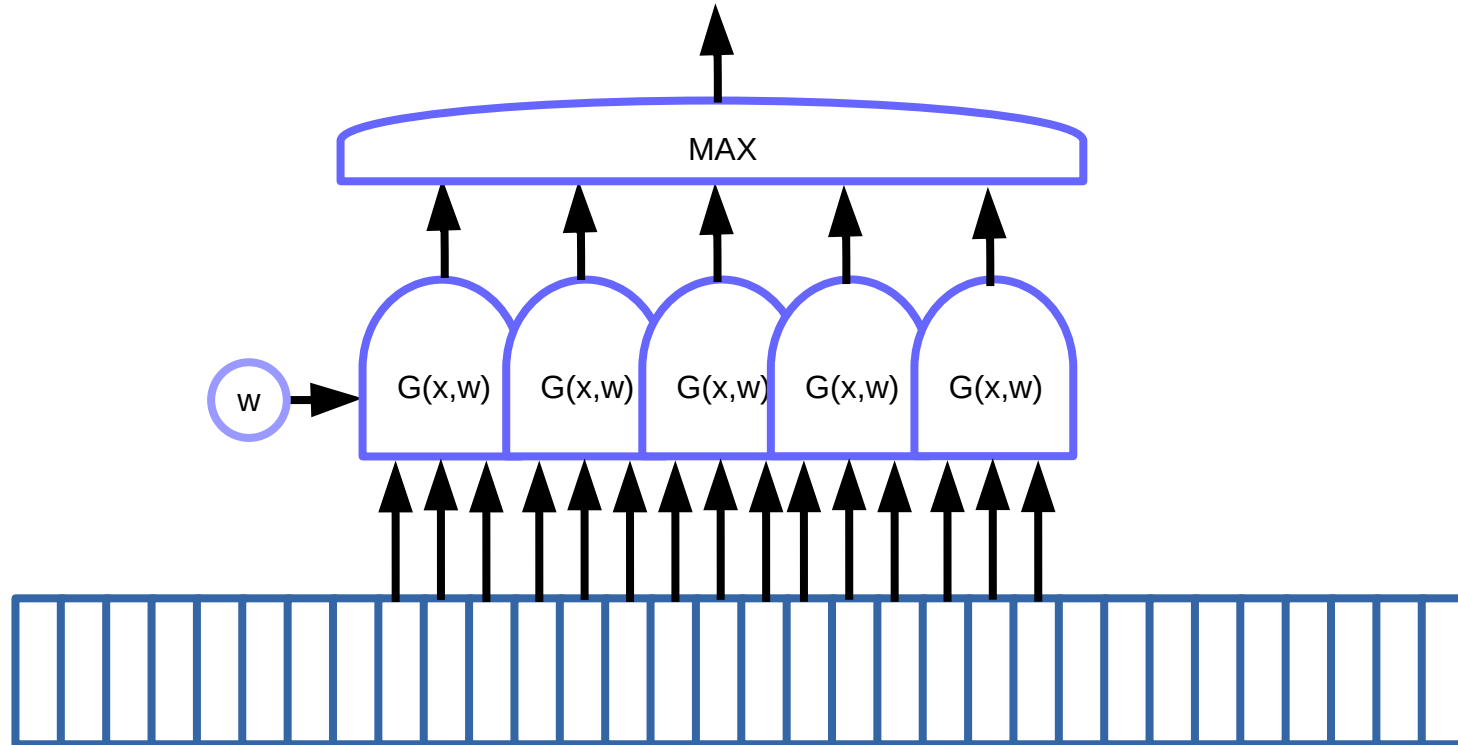
- ▶ Gradients are summed in the backprop

- ▶ The gradients w.r.t. shared parameters are added.



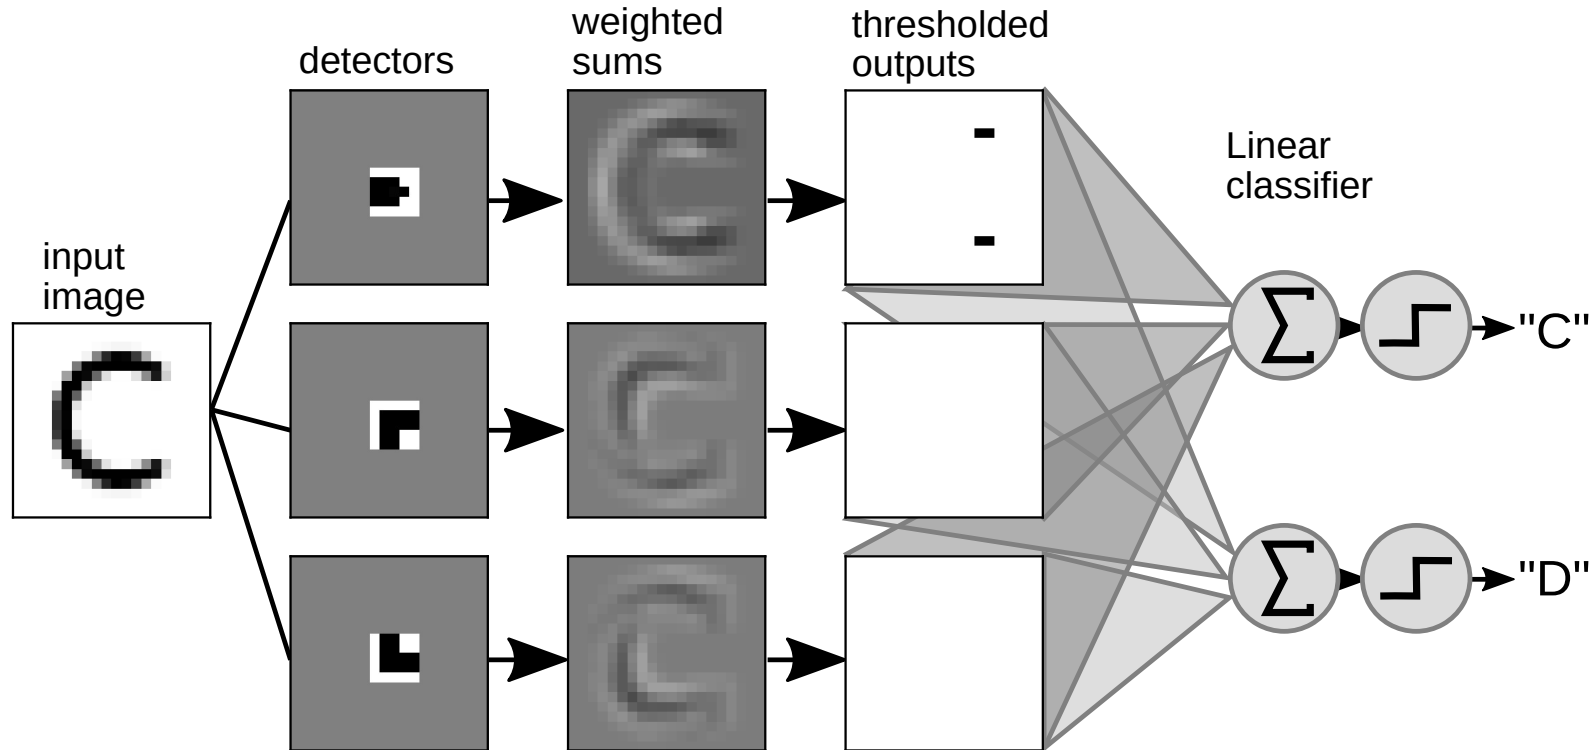
Shared Weights for Motif Detection

- ▶ Detecting motifs anywhere on an input



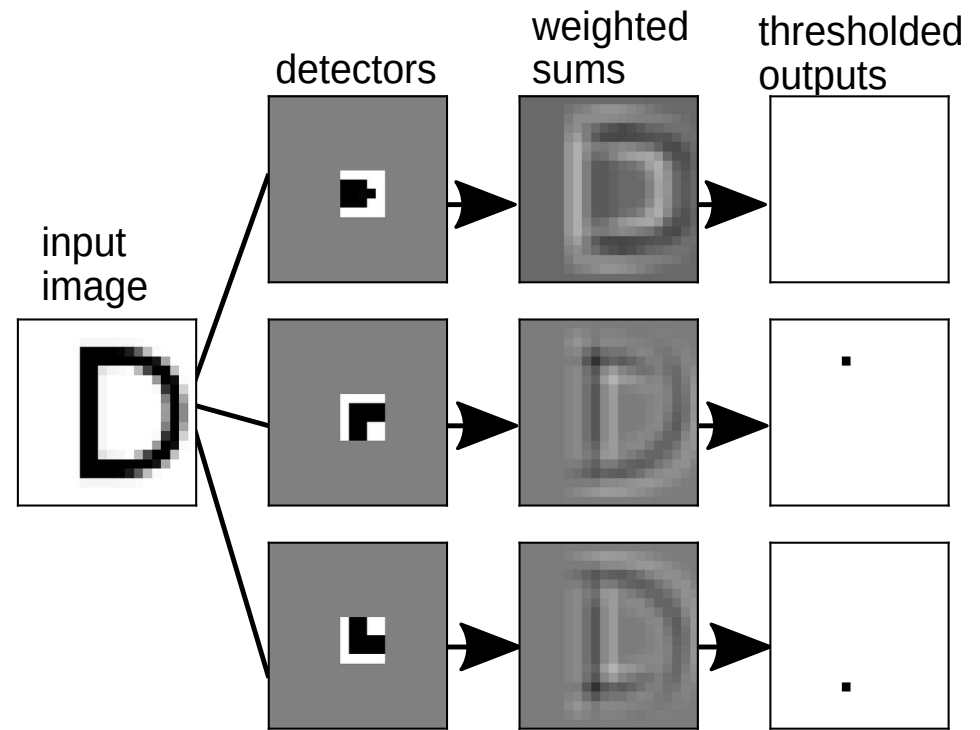
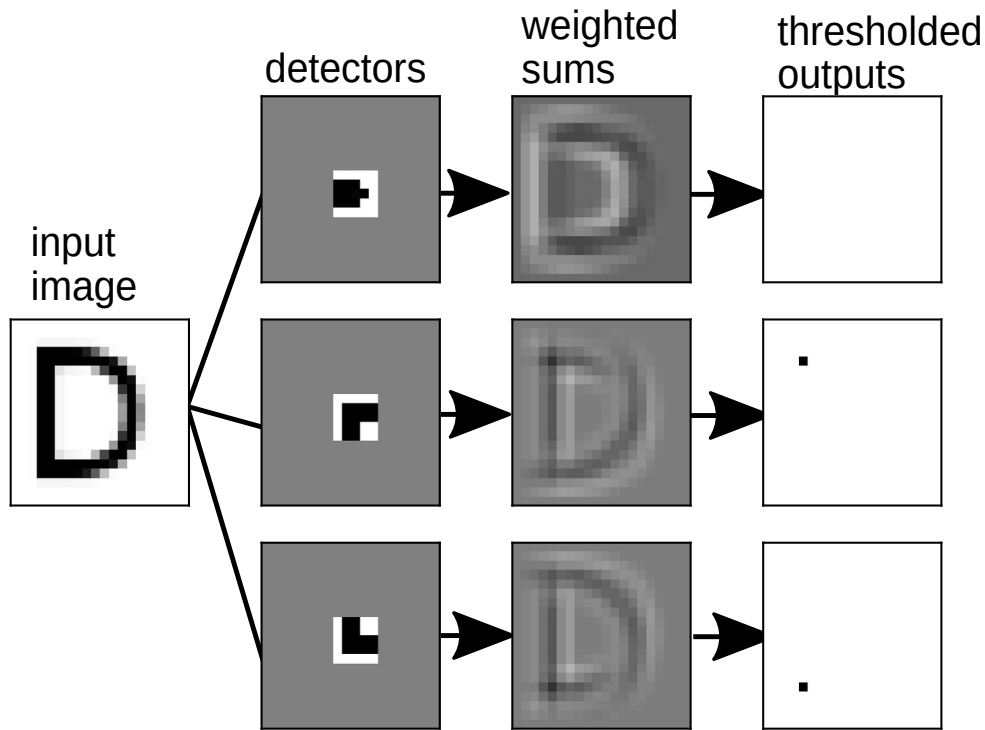
Detecting Motifs in Images

- Swipe “templates” over the image to detect motifs



Detecting Motifs in Images

► Shift invariance



Discrete Convolution (or cross-correlation)

- ▶ **Definition**
 - ▶ convolution

$$y_i = \sum_j w_j x_{i-j}$$

- ▶ **In practice**
 - ▶ Cross-correlation

$$y_i = \sum_j w_j x_{i+j}$$

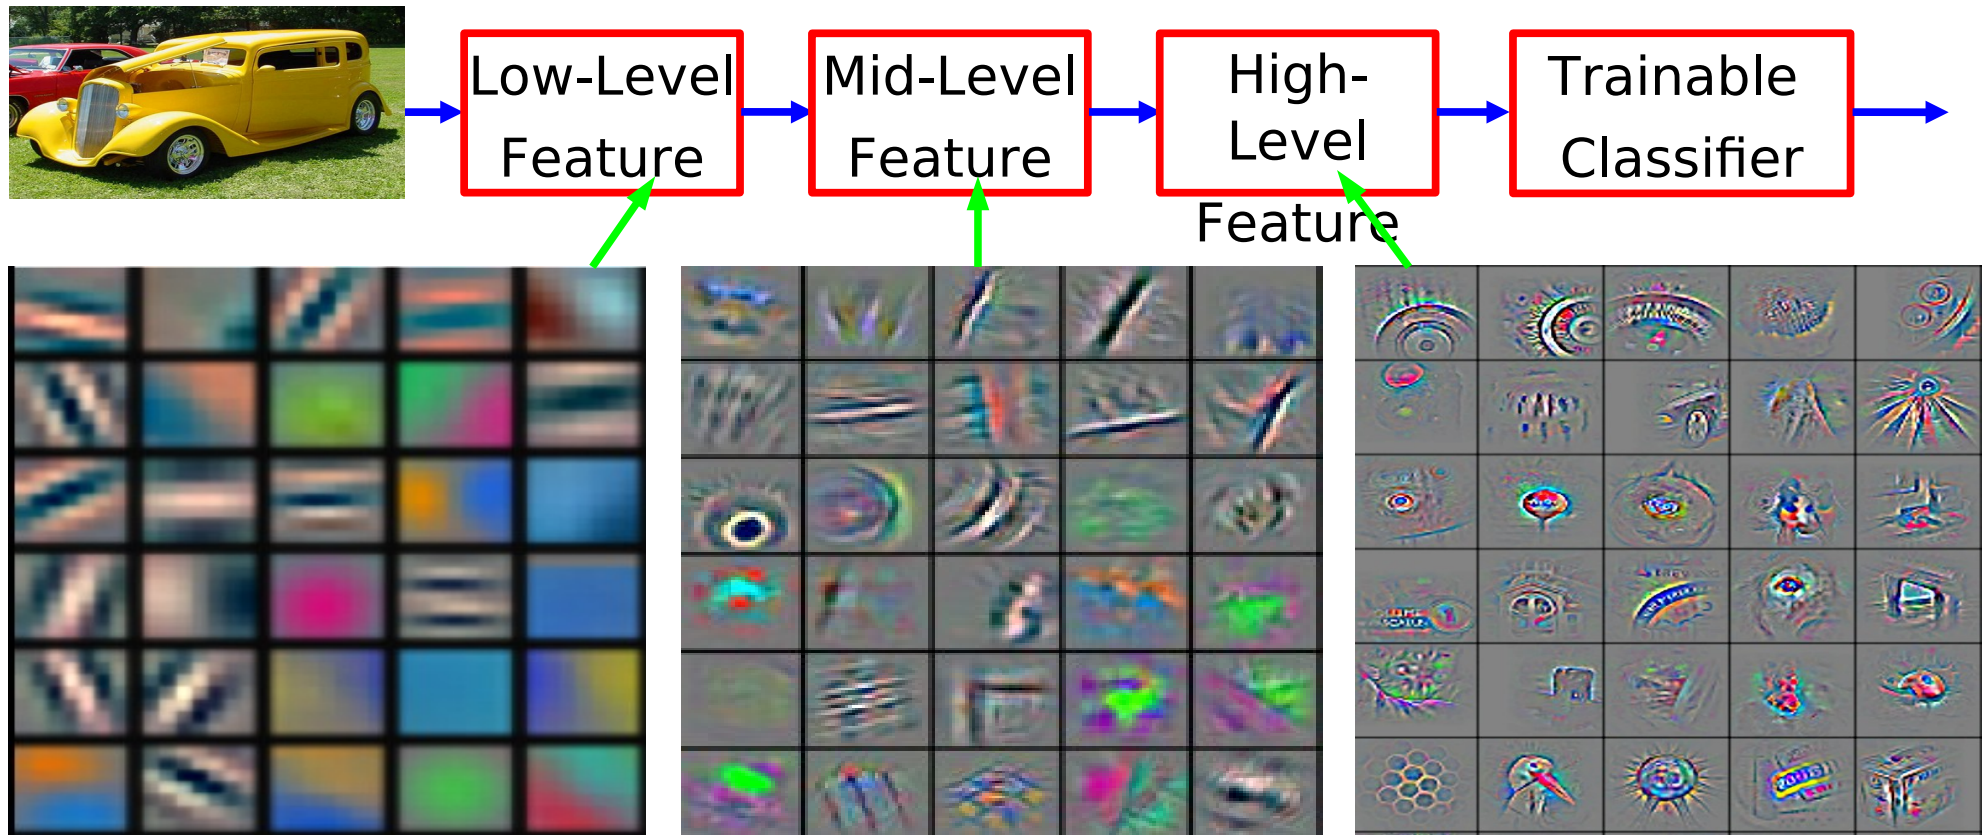
If we read w backwards, cross-correlation ends up to the equivalent to convolution.
Mathematically, convolution notation is more convenient for certain properties to have consistent form.
Programatically, w and x move on the same direction is more intuitive. DL frameworks choose the programatic way.

- ▶ **In 2D**

$$y_{ij} = \sum_{kl} w_{kl} x_{i+k, j+l}$$

Deep Learning = Learning Hierarchical Representations

It's **deep** if it has **more than one stage** of non-linear feature transformation

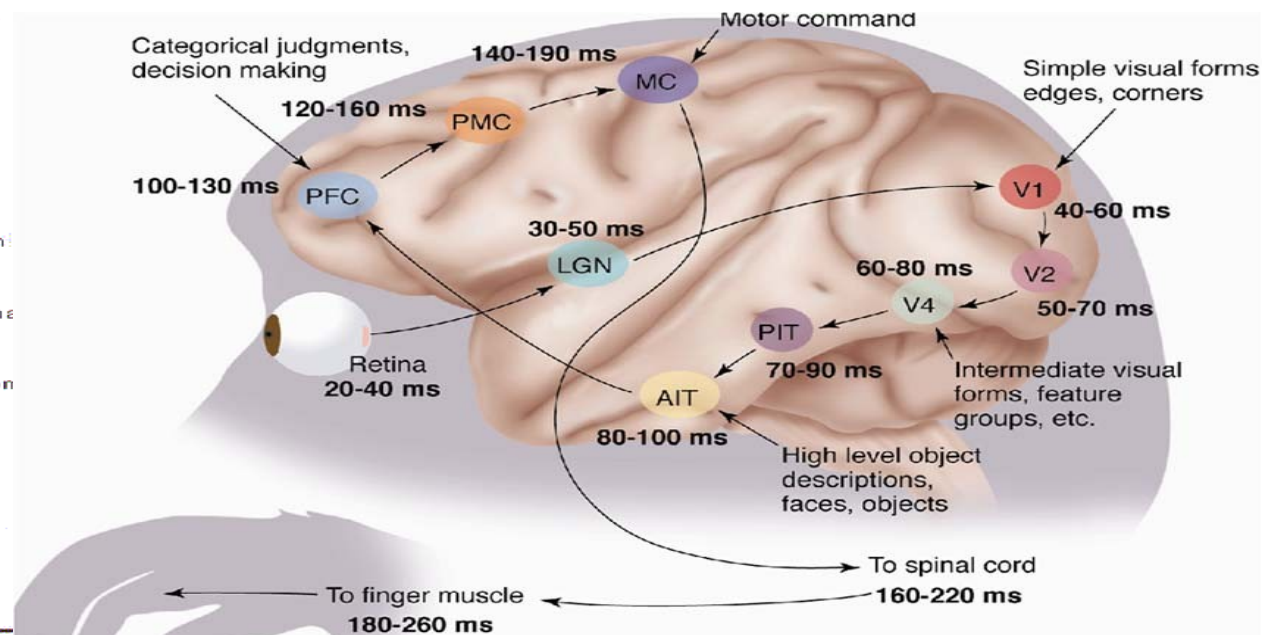
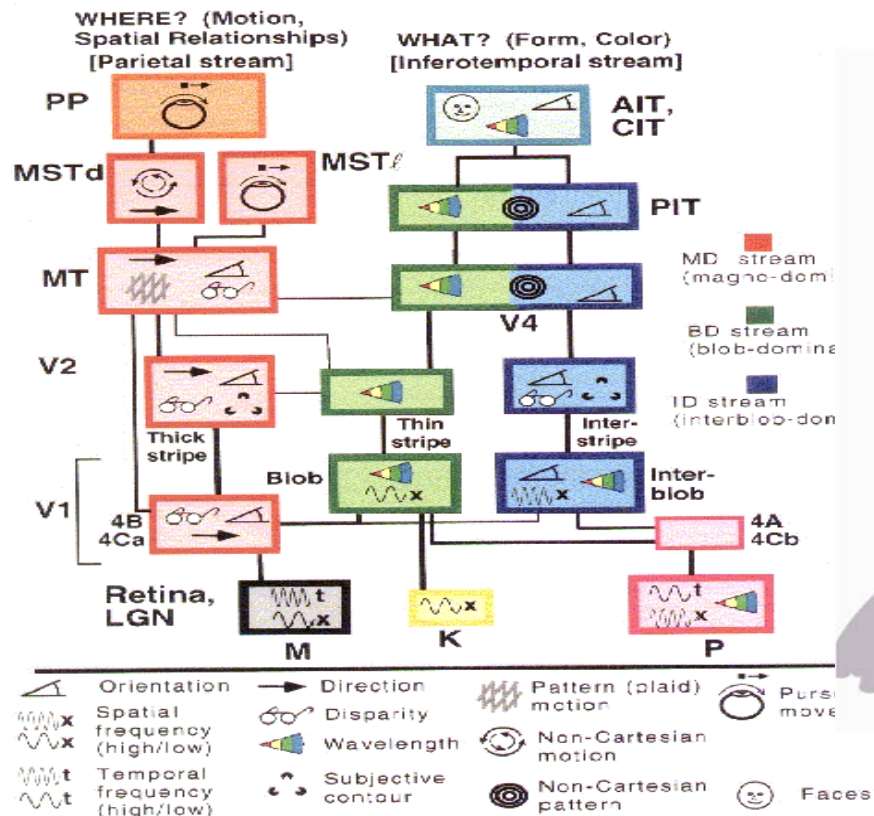


Feature visualization of convolutional net trained on ImageNet from [Zeiler & Fergus 2013]

How does the brain interpret images?

■ The ventral (recognition) pathway in the visual cortex has multiple stages

■ Retina - LGN - V1 - V2 - V4 - PIT - AIT



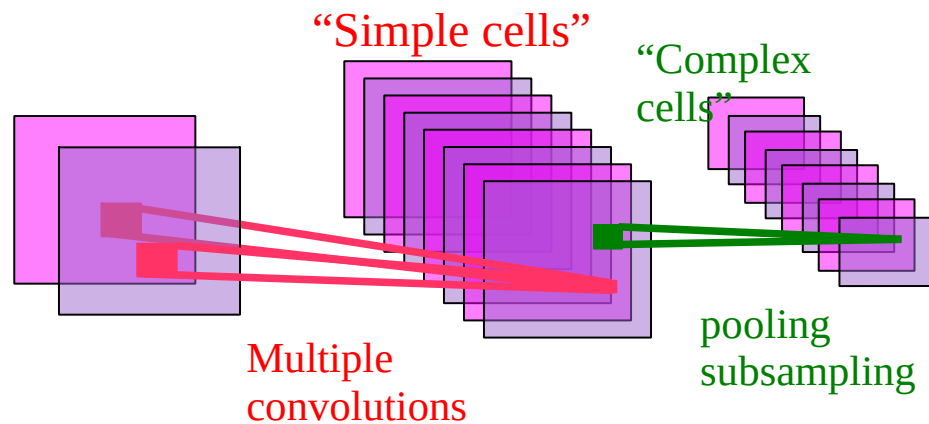
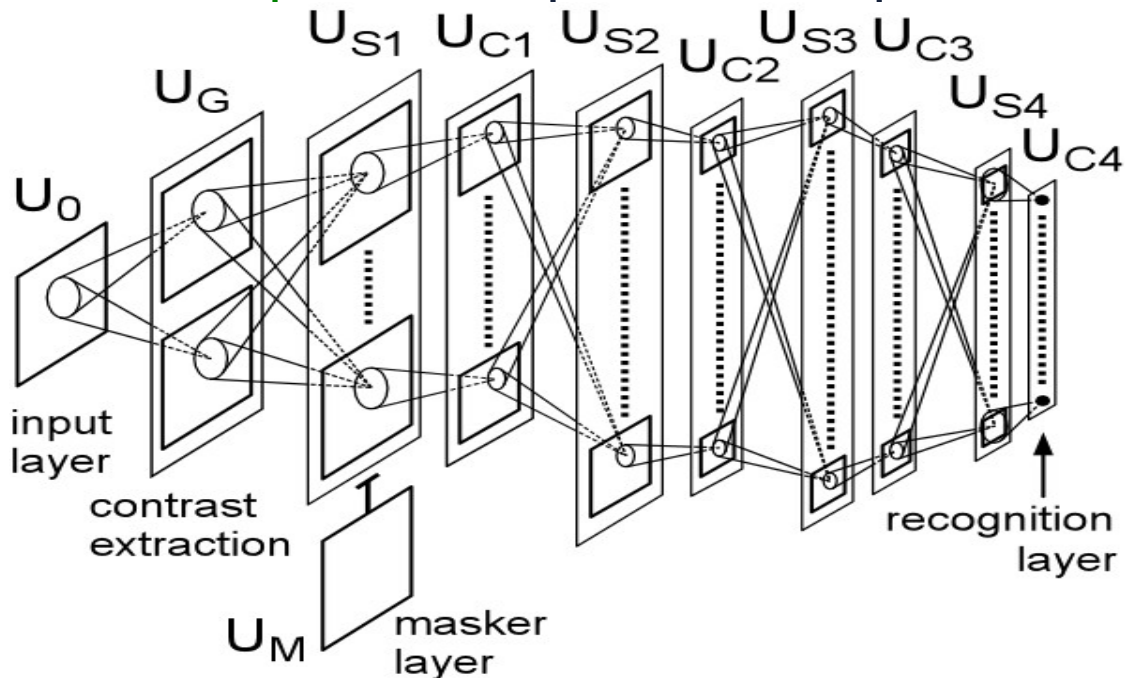
[picture from Simon Thorpe]

[Gallant & Van Essen]

Hubel & Wiesel's Model of the Architecture of the Visual Cortex

■ [Hubel & Wiesel 1962]:

- ▶ **simple cells** detect local features
- ▶ **complex cells** “pool” the outputs of simple

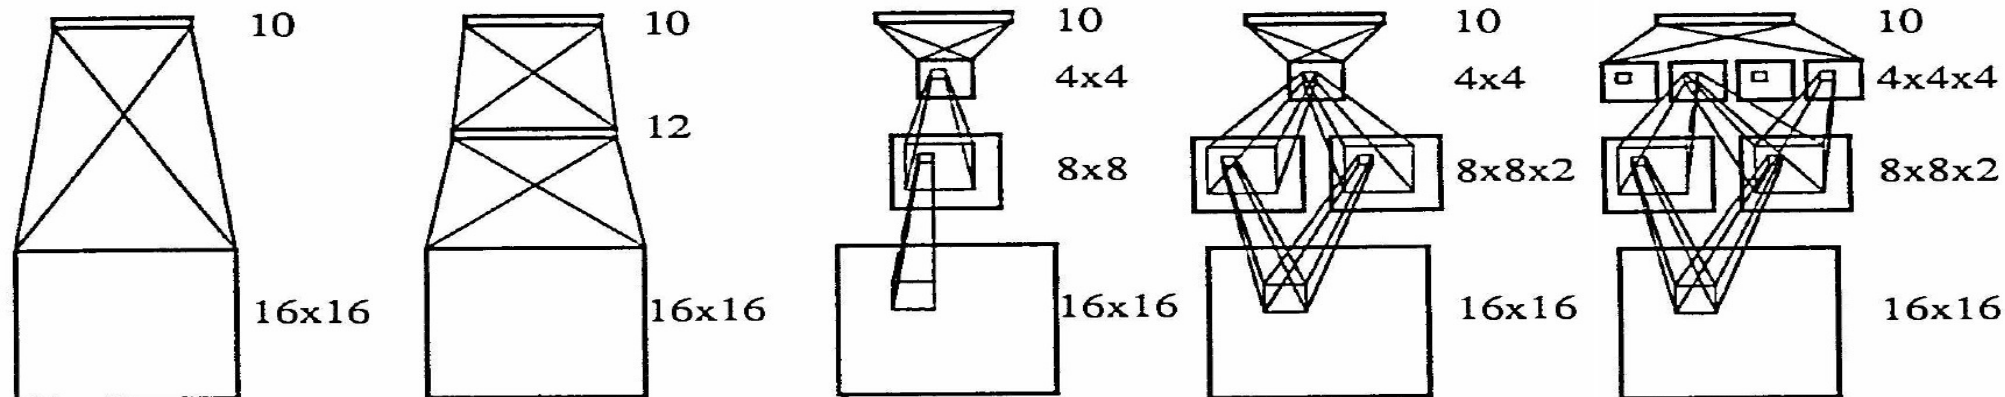


[Fukushima 1982][LeCun 1989, 1998],[Riesenhuber 1999].....

First ConvNets (U Toronto)[LeCun 88, 89]



► Trained with Backprop. 320 examples.



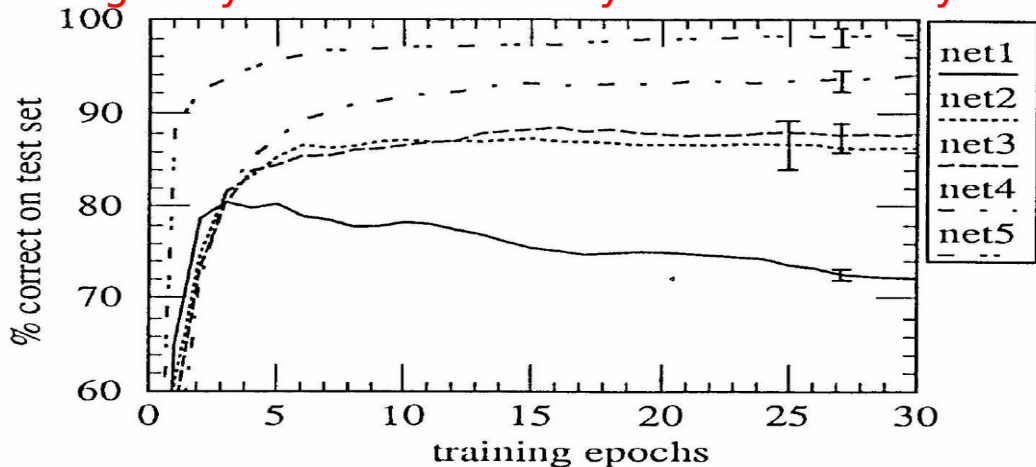
Single layer

Two layers FC

locally connected

Shared weights

Shared weights



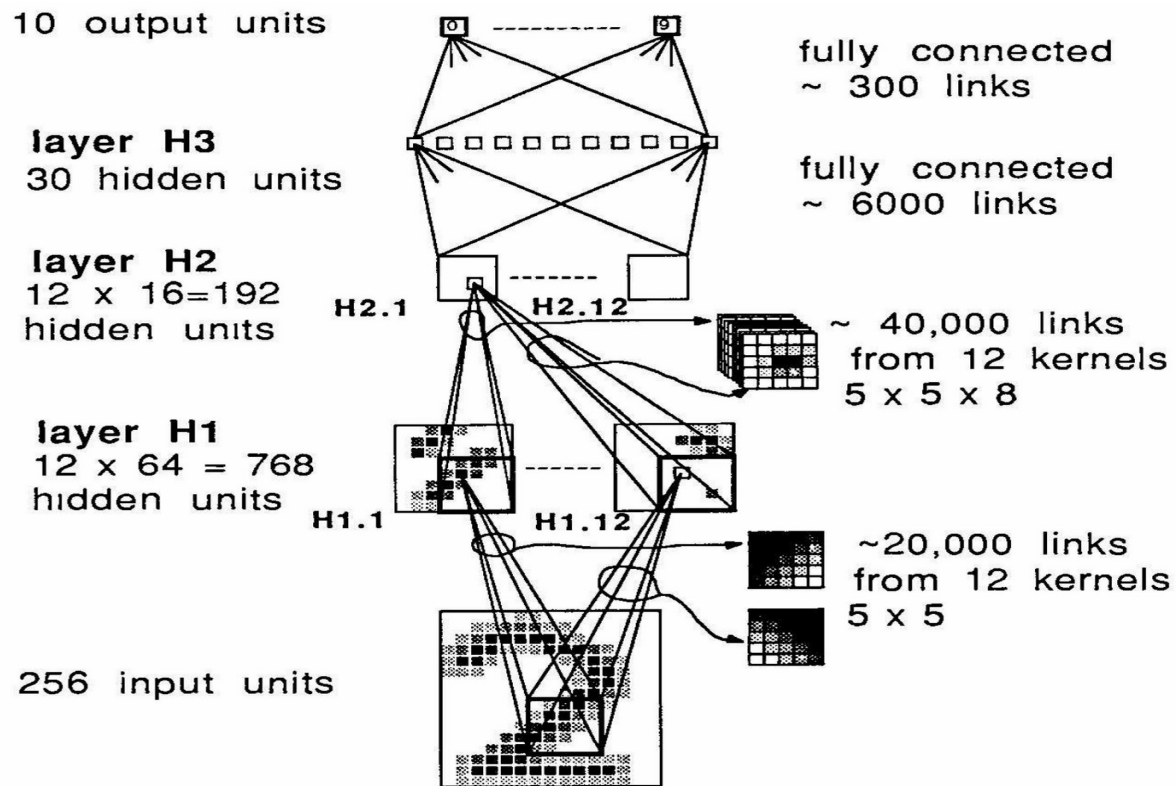
- Convolutions with stride
(subsampling)

- No separate pooling layers

network architecture	links	weights	performance
single layer network	2570	2570	80 %
two layer network	3240	3240	87 %
locally connected	1226	1226	88.5 %
constrained network	2266	1132	94 %
constrained network 2	5194	1060	98.4 %

First "Real" ConvNets at Bell Labs [LeCun et al 89]

- ▶ Trained with Backprop.
- ▶ USPS Zipcode digits: 7300 training, 2000 test
- ▶ Convolution with stride. No separate pooling.



80322-4129 80206

40004 14310

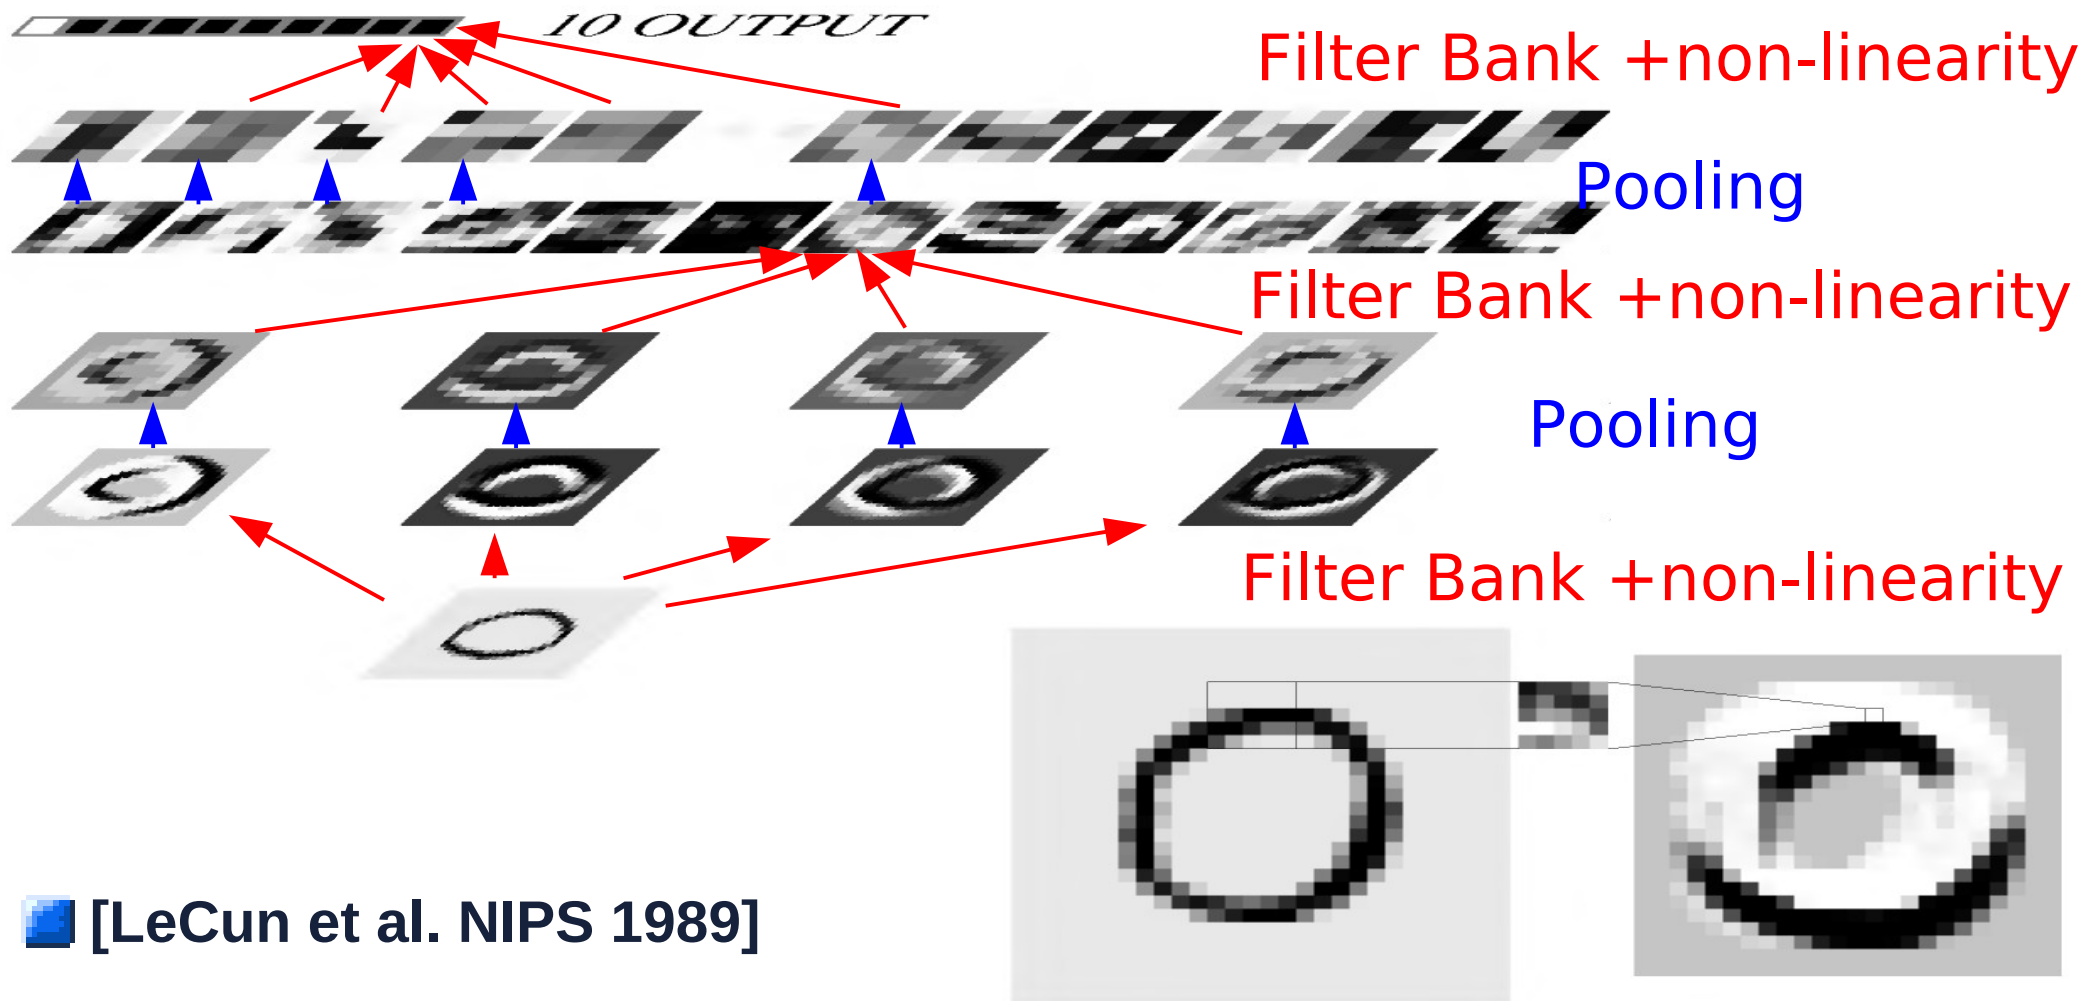
37872 05453

5502 75216

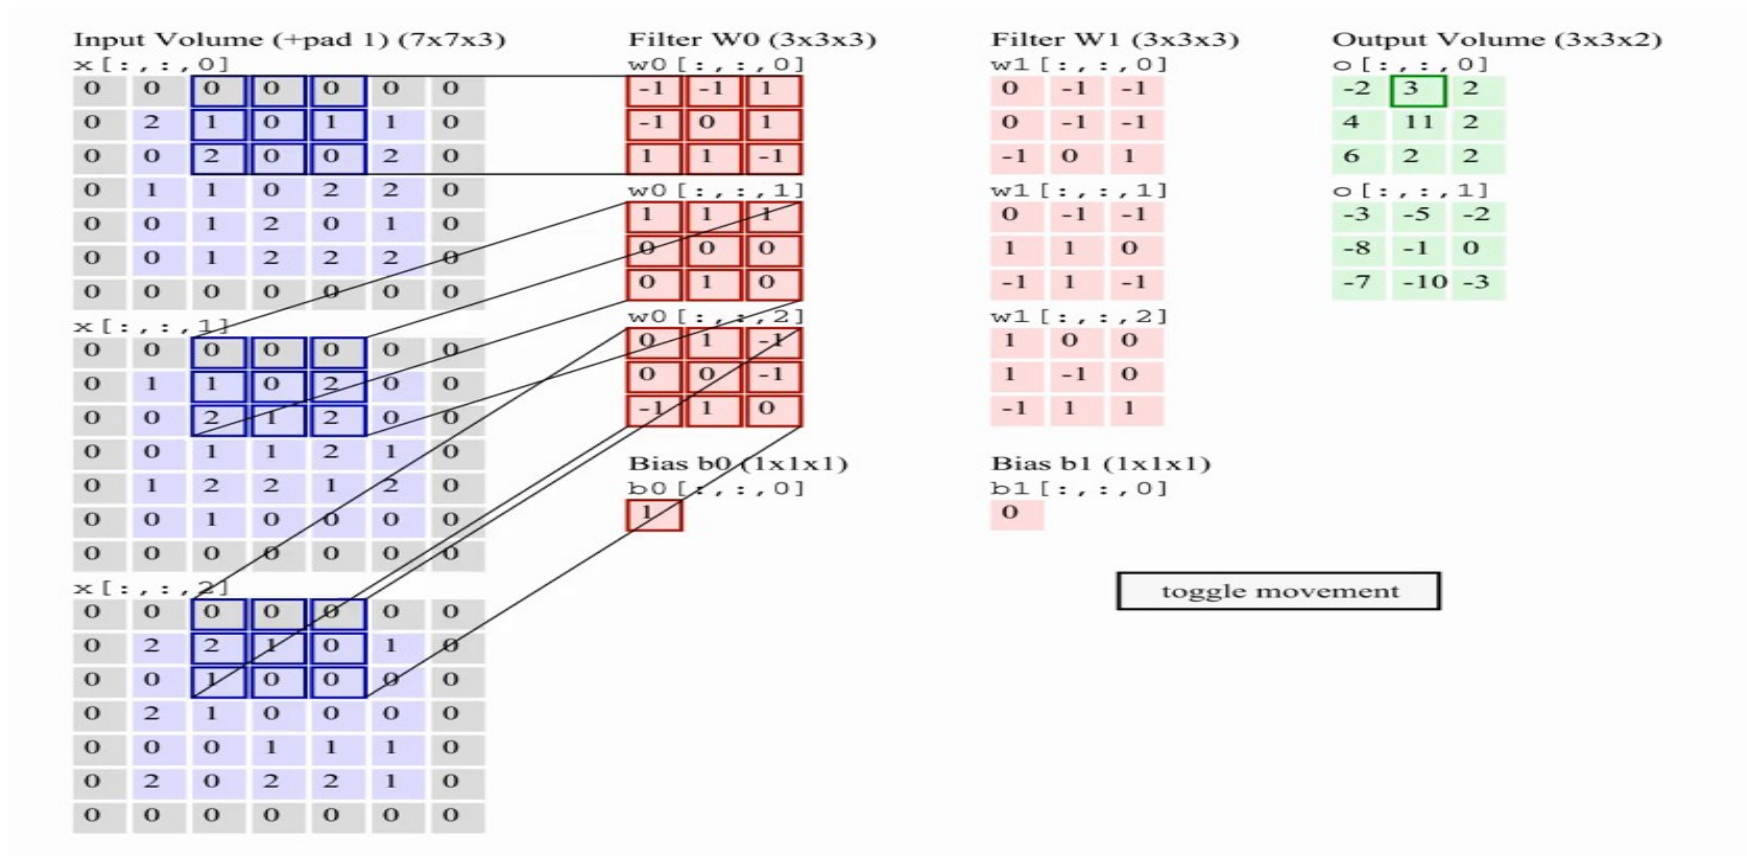
35460 44209

1611915485726803226414186
 6359720299299722510046701
 3084111591010615406103631
 1064111030475262009979966
 8912056708557131427955460
 2018750187112993089970984
 0109707597331972015519035
 1075318255182814358090943
 1787521655460554603546055
 18255108503047520439401

Convolutional Network Architecture



Multiple Convolutions



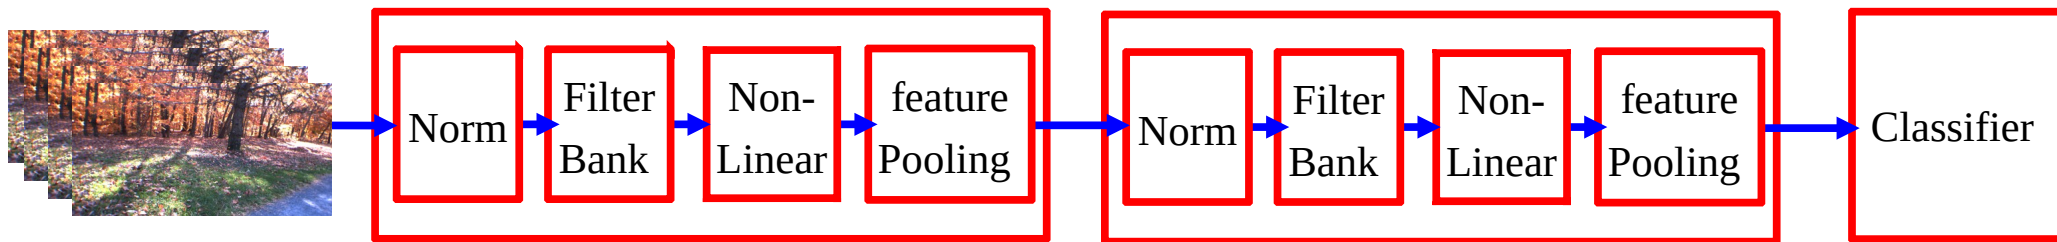
Animation: Andrej Karpathy <http://cs231n.github.io/convolutional-networks/>

Convolutional Network (vintage 1990)

Filters-tanh → pooling → filters-tanh → pooling → filters-tanh



Overall Architecture: multiple stages of Normalization → Filter Bank → Non-Linearity → Pooling



Normalization: variation on whitening (optional)

- Subtractive: average removal, high pass filtering
- Divisive: local contrast normalization, variance normalization

Filter Bank: dimension expansion, projection on overcomplete basis

Non-Linearity: sparsification, saturation, lateral inhibition....

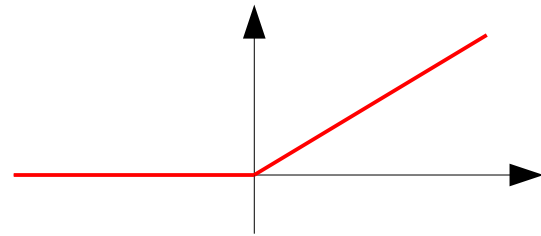
- Rectification (ReLU), Component-wise shrinkage, tanh,..

$$ReLU(x) = \max(x, 0)$$

Pooling: aggregation over space or feature type

- Max, Lp norm, log prob.

$$MAX : \max_i (X_i) ; \quad L_p : \sqrt[p]{X_i^p} ; \quad PROB : \frac{1}{b} \log \left(\sum_i e^{bX_i} \right)$$



LeNet5

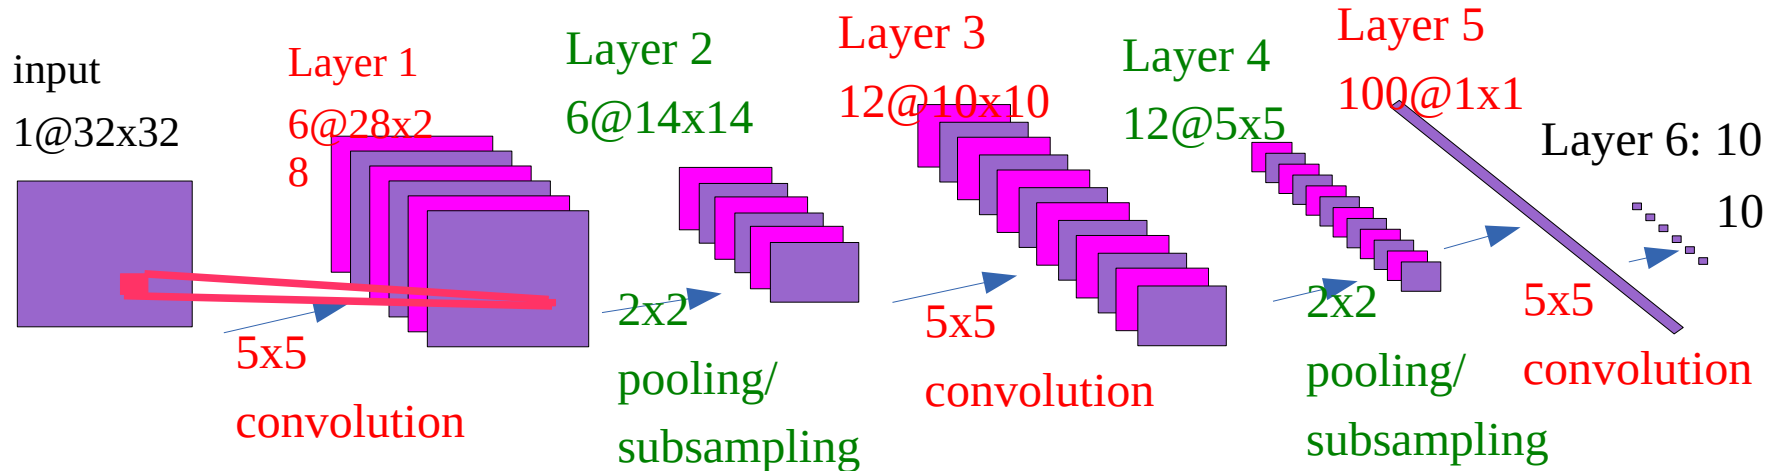
■ Simple ConvNet

■ for MNIST

■ [LeCun 1998]

■ PyTorch code ----->
- (slightly different net)

```
10 class Net(nn.Module):
11     def __init__(self):
12         super(Net, self).__init__()
13         self.conv1 = nn.Conv2d(1, 20, 5, 1)
14         self.conv2 = nn.Conv2d(20, 50, 5, 1)
15         self.fc1 = nn.Linear(4*4*50, 500)
16         self.fc2 = nn.Linear(500, 10)
17
18     def forward(self, x):
19         x = F.relu(self.conv1(x))
20         x = F.max_pool2d(x, 2, 2)
21         x = F.relu(self.conv2(x))
22         x = F.max_pool2d(x, 2, 2)
23         x = x.view(-1, 4*4*50)
24         x = F.relu(self.fc1(x))
25         x = self.fc2(x)
26         return F.log_softmax(x, dim=1)
```



LeNet5

■ Simple ConvNet

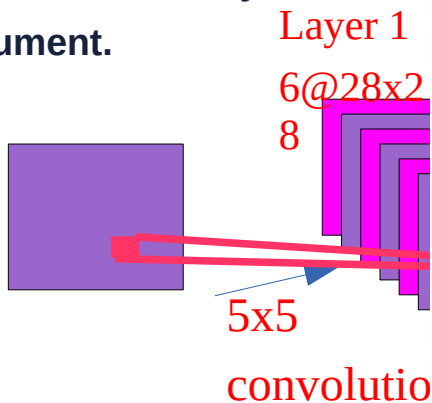
■ for MNIST

■ [LeCun 1998]

■ PyTorch code -- →

github.com/activatedgeek/LeNet-5

■ `nn.sequential()` with
ordered dictionary
argument.

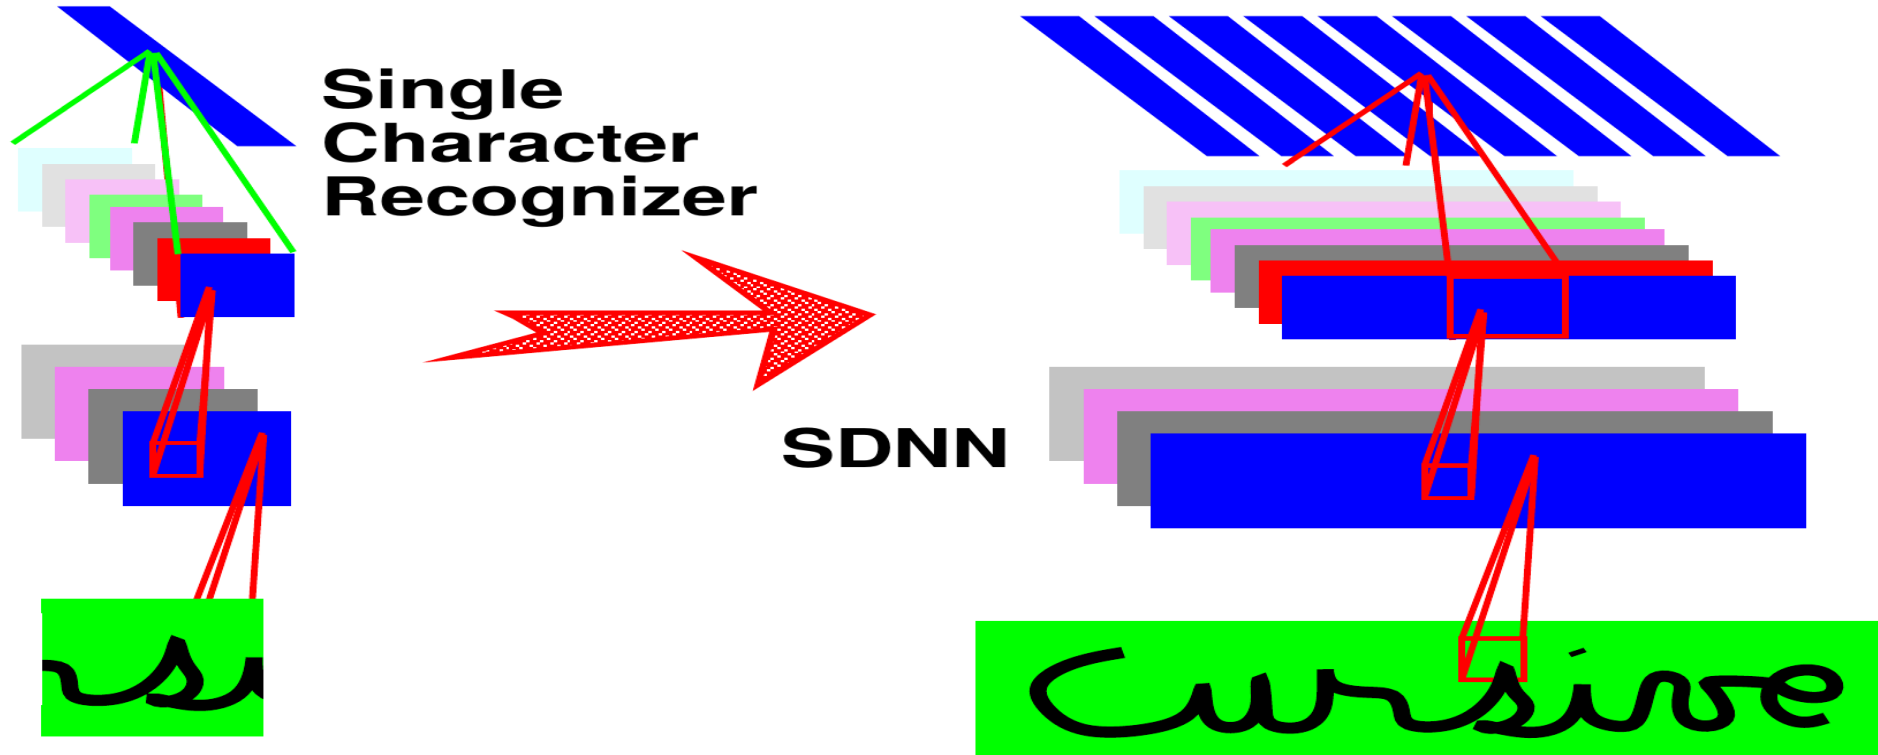


```
19 def __init__(self):
20     super(LeNet5, self).__init__()
21
22     self.convnet = nn.Sequential(OrderedDict([
23         ('c1', nn.Conv2d(1, 6, kernel_size=(5, 5))),
24         ('relu1', nn.ReLU()),
25         ('s2', nn.MaxPool2d(kernel_size=(2, 2), stride=2)),
26         ('c3', nn.Conv2d(6, 16, kernel_size=(5, 5))),
27         ('relu3', nn.ReLU()),
28         ('s4', nn.MaxPool2d(kernel_size=(2, 2), stride=2)),
29         ('c5', nn.Conv2d(16, 120, kernel_size=(5, 5))),
30         ('relu5', nn.ReLU())
31     ]))
32
33     self.fc = nn.Sequential(OrderedDict([
34         ('f6', nn.Linear(120, 84)),
35         ('relu6', nn.ReLU()),
36         ('f7', nn.Linear(84, 10)),
37         ('sig7', nn.LogSoftmax(dim=-1))
38     ]))
39
40 def forward(self, img):
41     output = self.convnet(img)
42     output = output.view(img.size(0), -1)
43     output = self.fc(output)
44     return output
```

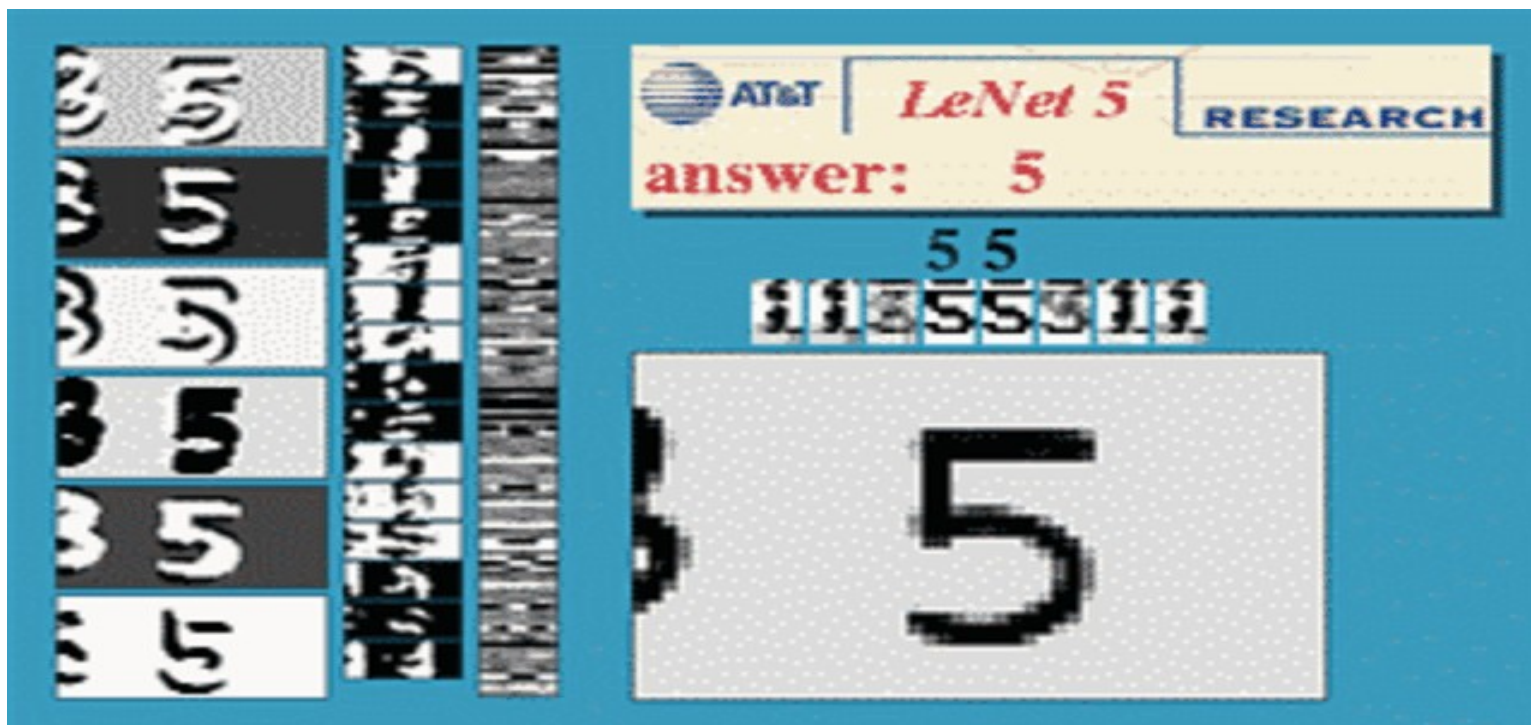
Multiple Character Recognition [Matan et al 1992]

■ Every layer is a convolution

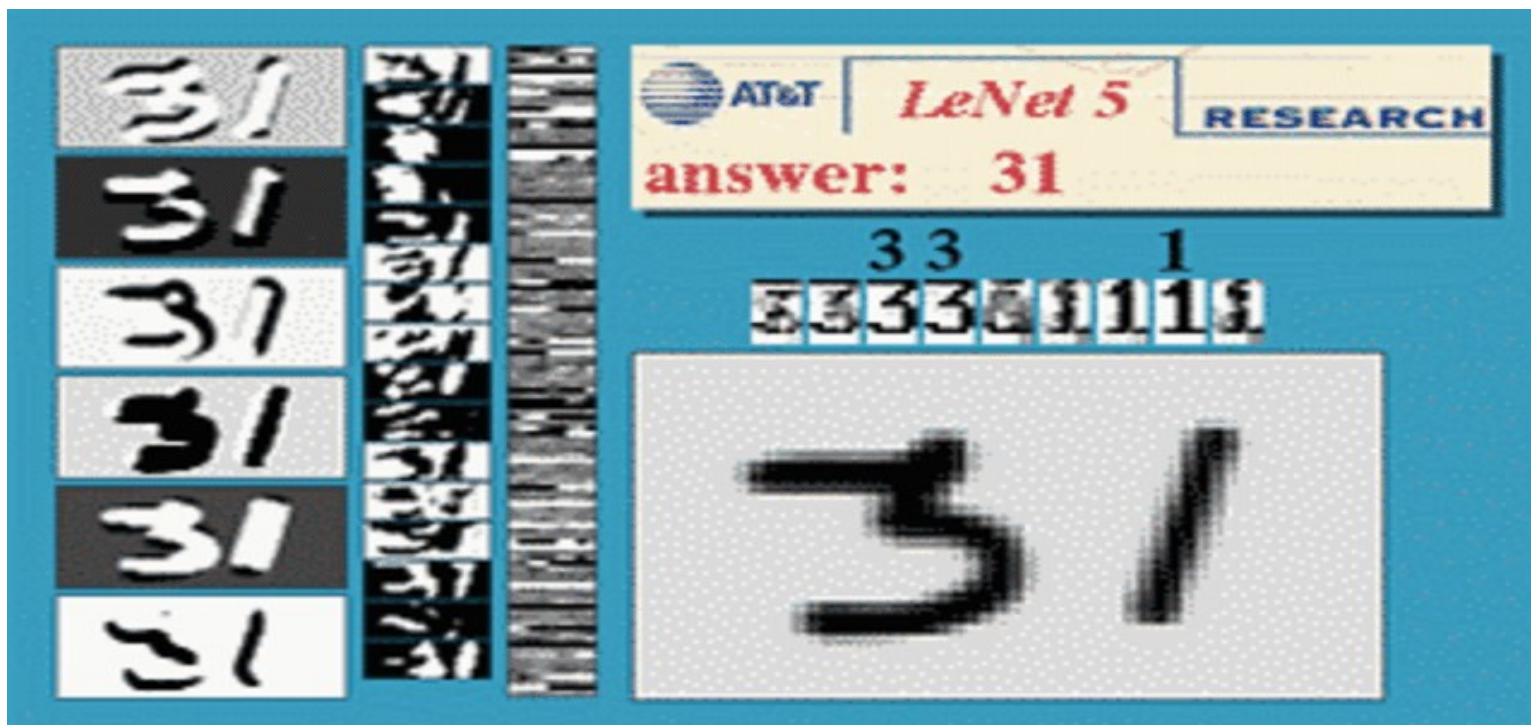
This is much cheaper than fixed size input and compute convolution again and again.
The last layer is 1x1 convolution



Sliding Window ConvNet + Weighted Finite-State Machine



Sliding Window ConvNet + Weighted FSM



What are ConvNets Good For

- Signals that comes to you in the form of (multidimensional) arrays.
- Signals that have strong local correlations
- Signals where features can appear anywhere
- Signals in which objects are invariant to translations and distortions.
- **1D ConvNets: sequential signals, text**
 - Text, music, audio, speech, time series.
- **2D ConvNets: images, time-frequency representations (speech and audio)**
 - Object detection, localization, recognition
- **3D ConvNets: video, volumetric images, tomography images**
 - Video recognition / understanding
 - Biomedical image analysis
 - Hyperspectral image analysis